

INTERNSHIP REPORT

DOMAIN: Machine Learning

TITLE: Restaurant Data Analysis

A Comprehensive Study of Predictive Analytics, Recommendation Systems,
and Location-Based Analysis in the Restaurant Industry

Submitted by: Rakshitha B N

**Internship Organization:
Cognifyz Technologies**

Project Tasks:

Task 1 - Restaurant Rating Prediction

Task 2 - Restaurant Recommendation System

Task 3 – Location Based Analysis

Executive Summary

This internship project at Cognifyz Technologies represents a comprehensive exploration of machine learning applications in the restaurant industry. Over the course of this project, I developed and implemented three major analytical systems that address real-world business challenges faced by restaurant stakeholders, from individual consumers to business owners and industry analysts.

The project encompasses three distinct but interconnected tasks. First, I built a Restaurant Rating Prediction system using multiple regression algorithms, achieving an impressive 96% accuracy with Random Forest. This predictive model can help restaurant owners understand which factors most strongly influence customer ratings, enabling data-driven business decisions. Second, I developed a Content-Based Restaurant Recommendation System that leverages natural language processing and cosine similarity to suggest restaurants matching user preferences. Finally, I conducted a comprehensive Location-Based Analysis that reveals geographical patterns in restaurant distribution, pricing, and customer preferences.

Through this project, I gained hands-on experience with the complete machine learning workflow: from data preprocessing and feature engineering to model training, evaluation, and deployment. The work not only strengthened my technical skills in Python, scikit-learn, and data visualization but also enhanced my ability to translate business problems into technical solutions. This report documents the methodologies, findings, challenges, and insights gained throughout the internship.

Table of Contents

1. Introduction.....	4
1.1 Project Background.....	4
1.2 Objectives and Scope.....	4
2. Literature Review and Background.....	5
3. Task 1: Restaurant Rating Prediction.....	6
4. Task 2: Restaurant Recommendation System.....	8
5. Task 4: Location-Based Analysis.....	10
6. Implementation and Technologies.....	11
7. Challenges and Solutions.....	12
8. Results and Key Insights.....	13
9. Future Work and Recommendations.....	14
10. Conclusion.....	15

1. Introduction

1.1 Project Background

The restaurant industry is a dynamic and competitive sector where data-driven insights can provide significant competitive advantages. With the proliferation of online review platforms and food delivery applications, restaurants now generate vast amounts of data about customer preferences, ratings, pricing strategies, and geographical performance. However, many businesses lack the analytical capabilities to extract actionable insights from this data.

This project was undertaken as part of my internship at Cognifyz Technologies, where I was tasked with developing machine learning solutions that address three key challenges in the restaurant domain: predicting customer ratings, recommending restaurants based on user preferences, and analyzing geographical distribution patterns. The dataset used contains comprehensive information about restaurants including cuisines offered, average ratings, price ranges, location coordinates, online delivery availability, and customer engagement metrics (votes).

1.2 Objectives and Scope

The primary objectives of this project were:

- To develop and evaluate multiple machine learning models for predicting restaurant ratings based on various features
- To build a recommendation system that suggests restaurants matching user preferences using content-based filtering
- To perform comprehensive location-based analysis revealing geographical patterns in restaurant distribution and performance
- To gain practical experience in the complete machine learning workflow including data preprocessing, feature engineering, model training, and evaluation

2. Literature Review and Background

2.1 Machine Learning in Restaurant Analytics

Machine learning has found extensive applications in the food and hospitality industry. Rating prediction, a form of regression analysis, helps businesses understand the factors that drive customer satisfaction. Research has shown that features such as cuisine type, price range, location, and service quality are strong predictors of customer ratings. In this project, I explored three regression algorithms—Linear Regression, Decision Tree, and Random Forest—to identify the most effective approach.

2.2 Recommendation Systems

Recommendation systems have become essential for online platforms, helping users discover relevant options from vast catalogs. Content-based filtering, which I implemented in this project, recommends items based on feature similarity rather than collaborative filtering (which relies on user behavior patterns). This approach is particularly suitable when user interaction data is limited or when cold-start problems need to be addressed. By using CountVectorizer and cosine similarity, the system can quantify how similar restaurants are based on their attributes.

2.3 Geographical Analysis in Business Intelligence

Location-based analysis provides valuable insights into market dynamics, competition density, and regional preferences. Geographical information systems (GIS) and interactive mapping tools like Folium enable stakeholders to visualize spatial patterns that would be difficult to discern from tabular data alone. Understanding where restaurants are concentrated, which areas have higher ratings, and how pricing varies by location can inform expansion strategies and market positioning.

3. Task 1: Restaurant Rating Prediction

3.1 Objective and Methodology

The goal of this task was to build a predictive model that estimates restaurant ratings based on various attributes. Accurately predicting ratings can help restaurant owners identify which aspects of their business most significantly impact customer perception. The methodology involved several key steps:

Data Preprocessing: The dataset required thorough cleaning, including handling missing values and encoding categorical variables. I used Label Encoding for binary features (such as online delivery availability) and One-Hot Encoding for multi-category features (such as cuisine types and price range). This preprocessing was crucial because machine learning algorithms require numerical inputs.

Feature Selection: I selected features that intuitively impact ratings: cuisines, location (city), price range, online delivery availability, and customer votes. The vote count turned out to be particularly influential, as it represents both popularity and customer engagement.

Model Training: I split the dataset into training (80%) and testing (20%) sets and trained three regression algorithms: Linear Regression (a baseline model), Decision Tree Regressor (capable of capturing nonlinear patterns), and Random Forest Regressor (an ensemble method that reduces overfitting).

3.2 Results and Analysis

The performance of each model was evaluated using R^2 score, which measures how well the model explains variance in the target variable:

Model	R^2 Score	Performance
Linear Regression	0.31	Weak
Decision Tree Regressor	0.92	Very Good
Random Forest Regressor	0.96	Excellent

Linear Regression performed poorly ($R^2 = 0.31$) because it assumes a linear relationship between features and ratings, which doesn't capture the complex, nonlinear patterns in restaurant data. Decision Tree achieved much better performance ($R^2 = 0.92$) by learning decision rules that split the data based on feature thresholds. Random Forest, an ensemble of multiple decision trees, achieved the best performance ($R^2 = 0.96$) by averaging predictions across trees, which reduces overfitting and improves generalization.

3.3 Key Insights

Feature importance analysis revealed that customer votes (engagement) were the most influential predictor of ratings. This makes an intuitive sense: restaurants that generate more customer engagement tend to have higher ratings because satisfied customers are more likely to leave positive reviews and recommend the restaurant. Price range and cuisine type also showed moderate importance, suggesting that customers have specific expectations based on these factors. Online delivery availability had a smaller but still measurable impact, indicating that convenience features contribute to overall satisfaction.

4. Task 2: Restaurant Recommendation System

4.1 Objective and Methodology

The objective of this task was to build a content-based recommendation system that suggests restaurants matching a user's preferences. Unlike collaborative filtering (which requires user interaction data), content-based filtering recommends items based solely on item features. This approach is ideal for scenarios where user history is unavailable or when addressing the cold-start problem.

The methodology involved:

Feature Engineering: I created a combined feature column by concatenating the restaurant's cuisines, city, price range, and online delivery status. For example, a restaurant might have the combined string 'North Indian New Delhi Medium Yes'. This textual representation captures the restaurant's key characteristics in a format suitable for natural language processing.

Vectorization: I used CountVectorizer from scikit-learn to convert these text descriptions into numerical vectors. CountVectorizer creates a vocabulary of all unique words across all restaurants and represents each restaurant as a vector counting how many times each word appears in its description.

Similarity Computation: Using cosine similarity, I computed the similarity between restaurants. Cosine similarity measures the angle between two vectors, with values ranging from 0 (completely different) to 1 (identical). Restaurants with similar cuisines, locations, and price ranges will have high cosine similarity scores.

4.2 Implementation and Results

The recommendation function takes user preferences (cuisine, city, price range, online delivery preference) as input, finds restaurants matching those criteria, and returns the top-N most similar restaurants. For example, when searching for 'North Indian restaurants in New Delhi with medium price range and online delivery', the system successfully recommended restaurants with matching characteristics.

One significant advantage of this content-based approach is its transparency: recommendations are based on explicit feature matches, making it easy to explain why a particular restaurant was suggested. Additionally, the system doesn't require historical user data, making it immediately useful for new users or restaurants.

4.3 Limitations and Future Enhancements

While effective, the current system has limitations. It only considers explicit features and cannot capture implicit preferences or discover unexpected options (the 'serendipity' problem). A hybrid approach combining content-based and collaborative filtering could address these limitations. Additionally, incorporating user reviews and sentiment analysis could provide richer recommendations by understanding qualitative aspects of customer experiences.

5. Task 4: Location-Based Analysis

5.1 Objective and Methodology

The objective of this task was to perform geographical and statistical analysis to understand how restaurant characteristics vary across locations. This analysis provides valuable insights for stakeholders considering market entry, expansion, or competitive positioning.

The analysis included several components:

- **Restaurant Distribution Analysis:** Examining which cities have the highest concentration of restaurants. New Delhi emerged as the city with the most restaurants, indicating a highly competitive market with diverse dining options.
- **Average Rating by City:** Comparing average ratings across cities to identify high-performing markets. Inner City showed the highest average ratings, suggesting either higher quality establishments or more discerning customer bases.
- **Price Range Analysis:** Understanding price distribution across locations to identify affordable vs. premium markets. This helps businesses understand local purchasing power and competition at different price points.
- **Cuisine Popularity Analysis:** Identifying which cuisines dominate different markets. North Indian cuisine was found to be the most common, reflecting regional culinary preferences.

5.2 Interactive Mapping

Using Folium, I created an interactive HTML map plotting restaurant location using latitude and longitude coordinates. Each marker on the map represents a restaurant and can be clicked to view details such as name, cuisine, rating, and price range. This visualization makes spatial patterns immediately apparent: clustering of restaurants in certain neighborhoods, gaps in coverage in other areas, and the relationship between location and ratings.

The interactive map serves multiple practical purposes: helping customers find nearby restaurants, enabling business owners to identify underserved areas, and assisting investors in evaluating

market opportunities. The ability to zoom, pan, and filter the map makes it a powerful exploratory tool.

6. Implementation and Technologies

The entire project was implemented in Python, leveraging its rich ecosystem of data science and machine learning libraries. Each technology was chosen for its specific strengths:

Pandas and NumPy: These libraries formed the foundation for data manipulation and numerical operations. Pandas provided powerful data structures (DataFrames) for handling tabular data, while NumPy enabled efficient array operations.

Scikit-learn: This comprehensive machine learning library provided implementations of regression algorithms, preprocessing utilities (Label Encoding, One-Hot Encoding), model evaluation metrics, and vectorization tools (CountVectorizer). Its consistent API made it easy to experiment with different algorithms.

Matplotlib and Seaborn: These visualization libraries enabled creation of informative plots including bar charts, histograms, and distribution plots. Seaborn's high-level interface made it particularly easy to create aesthetically pleasing visualizations with minimal code.

Folium: This library bridges Python and JavaScript visualization, creating interactive Leaflet maps from Python code. The resulting HTML maps could be easily shared and viewed in any web browser without requiring Python installation.

7. Challenges and Solutions

Throughout this project, I encountered several challenges that required creative problem-solving:

Challenge 1: Missing and Inconsistent Data

The dataset contained missing values and inconsistent formatting. For numerical features, I used median imputation to maintain distribution characteristics. For categorical features, I created an 'Unknown' category rather than dropping rows, preserving data volume for training. This approach balanced data quality with maintaining sufficient training samples.

Challenge 2: Encoding High-Cardinality Categorical Features

Features like 'cuisines' had many unique values, and simple one-hot encoding would create hundreds of columns. I addressed this by grouping less common cuisines into broader categories and keeping only the top N most frequent values. This reduced dimensionality while retaining the most important information.

Challenge 3: Balancing Model Complexity and Interpretability

While Random Forest achieved the best accuracy, its ensemble nature makes it harder to interpret than Decision Trees. To address this, I extracted feature importances from the Random Forest model, which helped identify which features most strongly influenced predictions. This allowed stakeholders to understand model behavior despite its complexity.

Challenge 4: Recommendation System Cold Start

Content-based filtering helped address the cold-start problem (recommending to new users without history), but it struggled to recommend restaurants with unique feature combinations. To mitigate this, I implemented a fallback mechanism that relaxes constraints progressively if no exact matches are found, ensuring users always receive recommendations.

8. Results and Key Insights

This project yielded several important insights with practical implications:

8.1 Business Insights

- **Customer Engagement is Critical:** The strong correlation between votes and ratings suggests that restaurants should focus on encouraging customer engagement through review platforms, loyalty programs, and social media presence.
- **Location Strategy Matters:** The concentration of restaurants in certain cities and the variation in average ratings by location indicate that market selection and positioning are crucial for success.
- **Cuisine Specialization:** North Indian cuisine dominance suggests established customer preferences but also presents opportunities for differentiation through less common cuisines.

8.2 Technical Insights

- **Ensemble Methods Excel:** Random Forest outperformed single models, demonstrating the value of ensemble approaches in capturing complex patterns.
- **Feature Engineering is Essential:** The success of the recommendation system depended heavily on effective feature combination and representation.
- **Visualization Enhances Understanding:** Interactive maps and statistical visualizations made patterns immediately apparent that would be difficult to identify in raw data.

9. Future Work and Recommendations

While this project successfully addressed its core objectives, several avenues for future enhancement exist:

Hybrid Recommendation System: Combining content-based filtering with collaborative filtering could improve recommendation quality by leveraging both item features and user behavior patterns. This would help address the serendipity problem and discover non-obvious recommendations.

Sentiment Analysis: Incorporating natural language processing to analyze customer reviews could provide richer insights into what drives ratings beyond numerical features. Understanding sentiment trends could help restaurants identify specific areas for improvement.

Deep Learning Models: Neural networks could potentially capture more complex nonlinear relationships, especially if more data becomes available. Embedding layers could better represent high-cardinality categorical features like cuisines and locations.

Real-Time Recommendation API: Deploying the recommendation system as a REST API would enable integration with web or mobile applications, making it accessible to end users in real-world scenarios.

Time-Series Analysis: If temporal data were available, analyzing how ratings and preferences change over time could reveal seasonal trends, identify rising cuisines, and predict future market dynamics.

10. Conclusion

This internship project at Cognifyz Technologies provided a comprehensive, hands-on experience in applying machine learning and data analysis to solve real-world business problems. Through three interconnected tasks—rating prediction, recommendation systems, and location-based analysis—I developed end-to-end solutions that address key challenges in the restaurant industry.

The rating prediction task demonstrated the power of ensemble methods, with Random Forest achieving 96% accuracy by effectively capturing complex patterns in customer behavior. The content-based recommendation system successfully matched user preferences to restaurant attributes, providing personalized suggestions without requiring historical interaction data. The geographical analysis revealed spatial patterns in restaurant distribution and performance, providing actionable insights for strategic decision-making.

Beyond technical skills, this project enhanced my ability to translate business requirements into technical solutions, handle real-world data challenges, and communicate findings effectively through visualization. The experience of working through the complete machine learning workflow—from data preprocessing and feature engineering to model evaluation and deployment—has been invaluable in developing practical expertise that complements theoretical knowledge.

The insights gained from this project have direct applications in the restaurant industry, from helping owners understand rating drivers to enabling platforms to provide better recommendations to users. More broadly, the methodologies developed here are transferable to other domains involving predictive analytics, recommendation systems, and geographical analysis.

I am grateful to Cognifyz Technologies for providing this opportunity to work on meaningful, industry-relevant projects. The internship has significantly strengthened my machine learning and data science capabilities, preparing me for future challenges in the field of artificial intelligence and analytics.

Submitted by: Rakshitha B N

May 2026